

Week 3 Cohort 4: R4DS Book Club

Chapters 4 & 5

Workflow: basics; Data transformation

Collin K. Berke

Twitter: @BerkeCollin

Last updated: 2020-12-30

5-minute ice breaker

- What's the last TV show/movie you watched?

Quick housekeeping/reminders

- Video camera is optional, but encouraged.
- If we need to slow down and discuss, let me know.
 - Most likely someone has the same question.
- Take time to learn the theory.
- Please attempt the chapter exercises.
- Please plan on teaching one of the lessons.

Tonight's discussion

- Chapter 4: Workflow basics
- Chapter 5: Data transformations
 - `filter()` rows
 - `arrange()` data
 - `select()` columns by their names
 - `mutate()` new variables
 - `group_by()` and `summarise()` a new calculation
 - `%>%` to chain commands together

Workflow basics

- R can be used like a calculator.

```
1 / 200 * 30
```

```
## [1] 0.15
```

```
(59 + 73 + 2) / 3
```

```
## [1] 44.66667
```

```
sin(pi / 2)
```

```
## [1] 1
```

Creating objects

- R is an object based programming language.
- Assign new objects with the `←` (Alt + -), avoid using `=`

```
# Avoid  
x = 3 * 4  
  
# object_name ← value  
# "Object name gets value"  
x ← 3 * 4
```

- What would happen if we type `x` in the terminal?

```
x  
#> [1] 12
```

- See all objects in the Environment pane

Naming rules

- Object names must start with a letter.

```
# Bad
```

```
2data
```

```
# Bad
```

```
_data
```

```
# Bad
```

```
%data
```

```
# Good
```

```
data
```

Naming rules (cont.)

- Only contain letters, numbers, `_`, and `.`

```
# Bad
```

```
flights_dep_&_arr
```

```
# Good
```

```
flights.filtered.data
```

```
# Good, though not very descriptive
```

```
flights_data_2
```


Naming rules (cont.)

- Use descriptive names.

```
# Bad
```

```
flights_2_data
```

```
# Good
```

```
flights_filtered_data
```

Naming rules (cont.)

- Choose a case and style, stay consistent.
 - Think about others reading/using your code.
 - Think about your future self reading/using your code.

```
# Examples from book
i_use_snake_case
otherPeopleUseCamelCase
some.people.use.periods
And_aFew.People_RENOUNCEconvention
```

Naming rules (cont.)

- Typos matter
- Case matters

"There's an implied contract between you and R: it will do the tedious computation for you, but in return, you must be completely precise in your instructions. Typos matter. Case matters. ~Hadley Wickham & Garrett Golemund (authors)"

```
y  
#> Error: object 'y' not found
```

- You won't always get an error from a typo, so keep an eye out.

```
# Functions used for importing data, we will discuss later  
read_csv() # from readr  
  
# vs.  
read.csv() # from utils
```

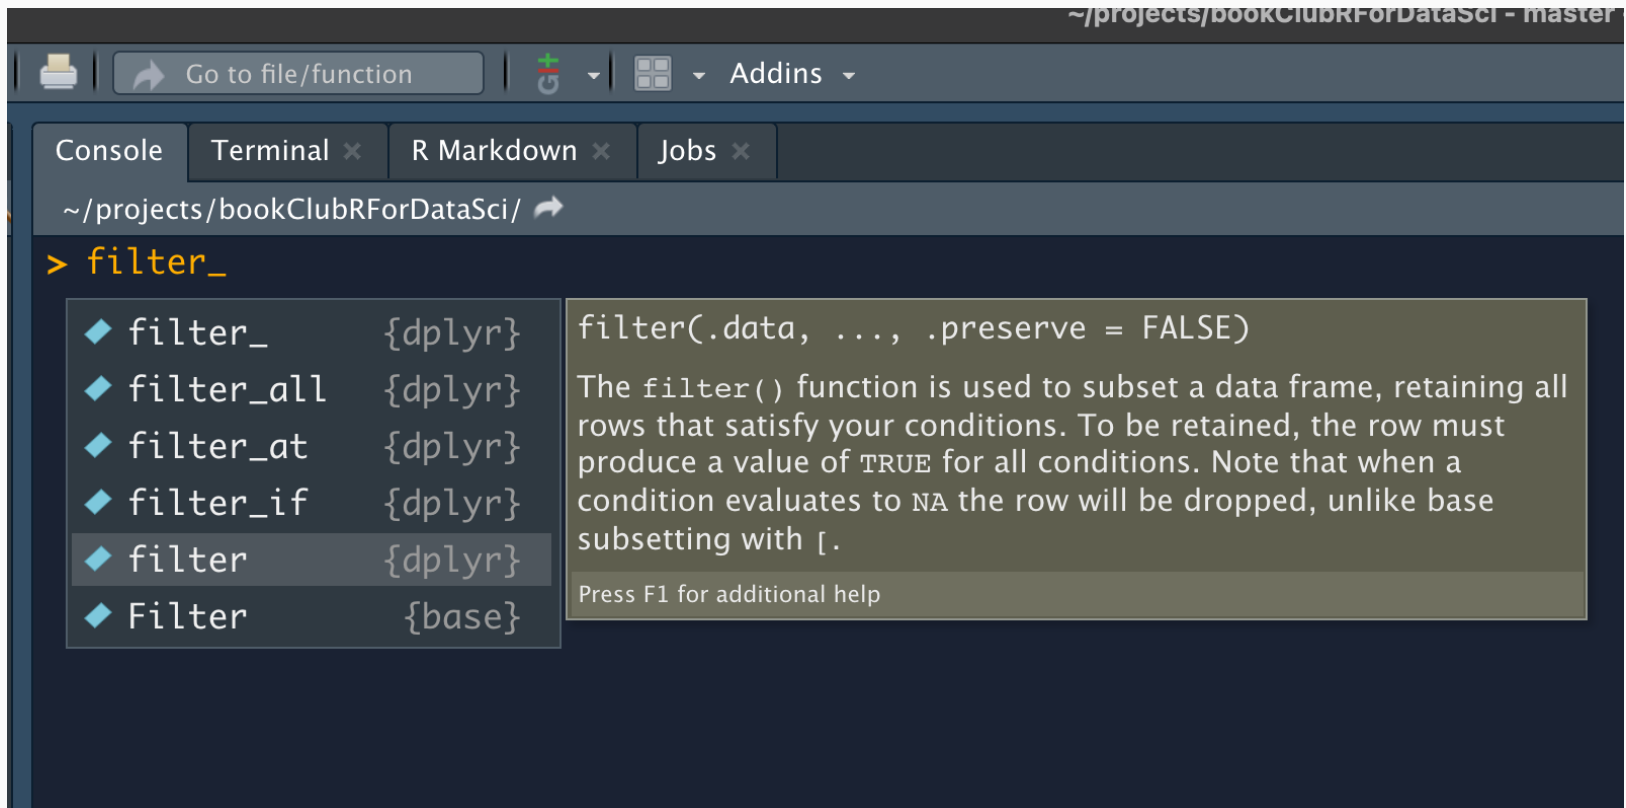
Naming rules, examples

tidyverse style guide

Google's R style guide

Calling functions

- Give your fingers a break, use the tab key.



The screenshot shows the RStudio interface with the console pane active. The path is `~/projects/bookClubRForDataSci/`. The prompt is `> filter_`. A list of functions is shown, with `filter` highlighted. A tooltip for `filter` is displayed, showing its signature and description.

Function	Package
<code>filter_</code>	<code>{dplyr}</code>
<code>filter_all</code>	<code>{dplyr}</code>
<code>filter_at</code>	<code>{dplyr}</code>
<code>filter_if</code>	<code>{dplyr}</code>
<code>filter</code>	<code>{dplyr}</code>
<code>Filter</code>	<code>{base}</code>

filter(.data, ..., .preserve = FALSE)

The `filter()` function is used to subset a data frame, retaining all rows that satisfy your conditions. To be retained, the row must produce a value of `TRUE` for all conditions. Note that when a condition evaluates to `NA` the row will be dropped, unlike base subsetting with `[]`.

Press F1 for additional help

Calling functions (cont.)

- Match `()` and `" "`.
 - If you see `+`, `R` is waiting for you to complete the expression.
 - Finish the expression or hit the `esc` key.

```
> library(nycflights13)
> filter(flights, carrier == "UA"
+
+ |
```

Chapter 5 - Data transformation

- **Note:** This chapter has lots of good information and many useful application examples. One session may not be enough to fully discuss each function and their many uses.

The verbs of data manipulation

Perform some action with our data

- `dplyr` functions:
 - `filter()`
 - `arrange()`
 - `select()`
 - `mutate()`
 - `group_by()` and `summarise()`

`dplyr` verbs never change our original data

- Get comfortable using the assignment operator, `←`.

The `nycflights13` data

```
library(nycflights13)
```

```
# more info, enter ?flights into console
```

```
names(flights)
```

```
##  [1] "year"           "month"          "day"            "dep_time"
##  [5] "sched_dep_time" "dep_delay"      "arr_time"       "sched_arr_t
##  [9] "arr_delay"      "carrier"        "flight"         "tailnum"
## [13] "origin"         "dest"           "air_time"       "distance"
## [17] "hour"           "minute"         "time_hour"
```

The `nycflights13` data types

```
# Quick view of the data types
glimpse(flights)
```

```
## Rows: 336,776
## Columns: 19
## $ year      <int> 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013, ...
## $ month     <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...
## $ day       <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...
## $ dep_time  <int> 517, 533, 542, 544, 554, 554, 555, 557, 557, 558, 558,...
## $ sched_dep_time <int> 515, 529, 540, 545, 600, 558, 600, 600, 600, 600, 600,...
## $ dep_delay <dbl> 2, 4, 2, -1, -6, -4, -5, -3, -3, -2, -2, -2, -2, -2, -...
## $ arr_time  <int> 830, 850, 923, 1004, 812, 740, 913, 709, 838, 753, 849...
## $ sched_arr_time <int> 819, 830, 850, 1022, 837, 728, 854, 723, 846, 745, 851...
## $ arr_delay <dbl> 11, 20, 33, -18, -25, 12, 19, -14, -8, 8, -2, -3, 7, -...
## $ carrier   <chr> "UA", "UA", "AA", "B6", "DL", "UA", "B6", "EV", "B6", ...
## $ flight    <int> 1545, 1714, 1141, 725, 461, 1696, 507, 5708, 79, 301, ...
## $ tailnum   <chr> "N14228", "N24211", "N619AA", "N804JB", "N668DN", "N39...
## $ origin    <chr> "EWR", "LGA", "JFK", "JFK", "LGA", "EWR", "EWR", "LGA"...
## $ dest      <chr> "IAH", "IAH", "MIA", "BQN", "ATL", "ORD", "FLL", "IAD"...
## $ air_time  <dbl> 227, 227, 160, 183, 116, 150, 158, 53, 140, 138, 149, ...
## $ distance  <dbl> 1400, 1416, 1089, 1576, 762, 719, 1065, 229, 944, 733,...
## $ hour      <dbl> 5, 5, 5, 5, 6, 5, 6, 6, 6, 6, 6, 6, 6, 6, 5, 6, 6, ...
## $ minute    <dbl> 15, 29, 40, 45, 0, 58, 0, 0, 0, 0, 0, 0, 0, 0, 59, ...
```

The `filter()` function

```
# data = data you want to filter
# filter exp = the expression used to filter data
filter(<data>, <filter exp>)
```

- Allows you to subset observations based on their values.
- Subset our data to return flights on Jan. 3, 2013

```
(jan1 ← filter(flights, month = 1, day = 1))
```

```
## # A tibble: 842 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517           515           2     830           819
## 2  2013     1     1     533           529           4     850           830
## 3  2013     1     1     542           540           2     923           850
## 4  2013     1     1     544           545          -1    1004          1022
## 5  2013     1     1     554           600          -6     812           837
## 6  2013     1     1     554           558          -4     740           728
## 7  2013     1     1     555           600          -5     913           854
## 8  2013     1     1     557           600          -3     709           723
## 9  2013     1     1     557           600          -3     838           846
```

The `filter()` function (cont.)

Effective filtering requires:

- The use of comparison operators.
 - `'>'` greater than
 - `'>='` greater than or equal to
 - `'<'` less than
 - `'<='` less than or equal to
 - `'!= '` not equal
 - `'=='` equal
- The application of logical operators.
 - `'&'` is "and"
 - `'|'` is "or"
 - `'!'` is "not"
- The use of the `%in%` operator

The `filter()` function, examples

- Using the `flights` data (you give it a try):
 - How many flights were to the major airports in Chicago (i.e., ORD, MDW) in 2013?
 - *There are multiple ways to get the answer.*

```
# One solution
filter(flights, dest == "ORD" | dest == "MDW")

# Another solution
filter(flights, dest %in% c("ORD", "MDW"))
```

The `arrange()` function

```
# data = data you want to use  
# col = column name you want to arrange by  
arrange(<data>, <col>, desc(<col>))
```

- `arrange()` - changes the order of the rows based on a variable(s).
- Orders rows in ascending (default) or descending (`desc()`) order.
- Multiple variables are used as tie-breakers.
- `NA` values are always sorted at the end.

The `arrange()` function, examples

- What will the result of this code be?

```
arrange(flights, desc(dep_delay))
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     9     641           900         1301    1242         1530
## 2  2013     6    15    1432          1935         1137    1607         2120
## 3  2013     1    10    1121          1635         1126    1239         1810
## 4  2013     9    20    1139          1845         1014    1457         2210
## 5  2013     7    22     845          1600         1005    1044         1815
## 6  2013     4    10    1100          1900          960    1342         2211
## 7  2013     3    17    2321           810          911     135         1020
## 8  2013     6    27     959          1900          899    1236         2226
## 9  2013     7    22    2257           759          898     121         1026
## 10 2013    12     5     756          1700          896    1058         2020
## # ... with 336,766 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

The `arrange()` function, examples (cont.)

- What will be the result of this code be?

```
filter(flights, is.na(dep_time))
```

```
## # A tibble: 8,255 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>         <int>
## 1  2013     1     1      NA           1630         NA      NA           1815
## 2  2013     1     1      NA           1935         NA      NA           2240
## 3  2013     1     1      NA           1500         NA      NA           1825
## 4  2013     1     1      NA            600         NA      NA            901
## 5  2013     1     2      NA           1540         NA      NA           1747
## 6  2013     1     2      NA           1620         NA      NA           1746
## 7  2013     1     2      NA           1355         NA      NA           1459
## 8  2013     1     2      NA           1420         NA      NA           1644
## 9  2013     1     2      NA           1321         NA      NA           1536
## 10 2013     1     2      NA           1545         NA      NA           1910
## # ... with 8,245 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```


The `select()` function

`select()` subsets the data with specific operations based on the names of the variables.

```
# data = data you want to use  
# col = column(s) you want to include in your selection  
select(<data>, <col>, <col>)
```

- Function is useful when you have datasets with many variables.

The `select()` function (cont.)

- `select()` can be combined with **helper functions** to select any combination of columns.
- `starts_with()` - matches names from the start.
- `ends_with()` - matches names from the end.
- `contains()` - matches names anywhere.
- `matches()` - match based on a regular expression.
- `num_range()` - matches based on a sequence of numbers.
- `everything()` - select all non-specified variables.

General applications of `select()`

```
# Select by name
select(flights, year, month, day)

# Select a range of columns (inclusive)
select(flights, year:day)

# Select columns excluding a range (inclusive)
select(flights, -(year:day))
```

Combining helper functions and `select()`

5.4.1.1. Brainstorm as many ways as possible to select `dep_time`, `dep_delay`, `arr_time`, and `arr_delay` from `flights`? (Anyone want to share?)

```
select(flights, dep_time, dep_delay, arr_time, arr_delay)

select(flights, ends_with(c("time", "delay")))

select(flights, starts_with(c("dep", "arr")))

select(flights, matches(c("^dep", "^arr")))

# Works, but it is ambiguous
select(flights, 4, 6, 7, 9)

select(flights, 4, 6:7, 9)

vars ← c("dep_time", "dep_delay", "arr_time", "arr_delay")

select(flights, all_of(vars))
```

Renaming and the use of `select()`

- Renaming can be done using `select`.
- However, the book suggests using `rename()` instead.

```
# Using rename(), keeps all the columns
# data = data you want to use
# new name = the new name you want to apply to the column
# old name = the old name of the column you want to change
rename(<data>, <new name> = <old name>)

rename(flights, tail_num = tailnum)

# Using select to rename and retain the one column
select(flights, tail_num = tailnum)
```

The `mutate()` function

`mutate()` adds new columns that are functions of existing columns.

- Always adds new columns at the end of the dataset.
- Only want the calculated columns?
 - Use `transmute()`.
- Utilize creation functions to calculate new columns.
 - **Must be vectorised** (i.e., takes a vector as input, returns a vector with the same number of values as output)

'Useful' creation functions

- Arithmetic operators
 - `+`, `-`, `*`, `/`, `^`
- Modular arithmetic
 - `%/%`, `%%`
- Logs
 - `log()`, `log2()`, `log10()`
- Offsets
 - `lead()`, `lag()`
- Cumulative and rolling aggregates
 - `cumsum()`, `cumprod()`, etc.
- Logical comparisons
- Ranking
 - `min_rank()`, `percent_rank()`
- **Note:** not an exhaustive list.

mutate() function, example

- What will be the result of the following code?

```
# What is happening here?
flights_sml <- select(flights,
  year:day,
  ends_with("delay"),
  distance,
  air_time
)

# What is the result?
mutate(flights_sml,
  gain = dep_delay - arr_delay,
  hours = air_time / 60,
  gain_per_hour = gain / hours # What concept is happening here?
)
```

- Notice how we can refer to previously created variables in the same `mutate()` function.

The `%>%` operator

The `%>%` (pipe) allows us to chain commands together.

As suggested by this reading, a good way to pronounce `%>%` when reading code is "then". ~Hadley Wickham & Garrett Grolemund (authors).

Why should we use the %>%?

```
# filter data for flights into Atlanta
flights_dest_atl <- filter(flights, dest = "ATL")

# Which flights had the longest arrival delay?
flights_atl_arr_delay <- arrange(flights_dest_atl, desc(arr_delay))

# I only care about carriers and arrival delays, use select
flights_carrier_delay <- select(flights_atl_arr_delay, carrier, arr_delay, arr_time)
```

```
# Same result, but with the `%>%`
# Use ctrl + shift + m on keyboard
flights_carrier_delay <- flights %>%
  filter(dest = "ATL") %>%
  arrange(desc(arr_delay)) %>%
  select(arr_time, arr_delay, carrier)
```

- It makes the code easier to read.
- Keeps us from cluttering up our environment with objects.
- What other benefits does the pipe provide?

The `summarise()` function

`summarise()` collapses a data frame to a single row.

```
# data = data you want to create a new variable from  
# new var name = the name your new variable will be  
# calc exp = the calculation you want to perform  
summarise(<data>, <new var name> = <calc exp>)
```

```
summarise(flights, delay = mean(dep_delay, na.rm = TRUE))
```

```
## # A tibble: 1 x 1  
##   delay  
##   <dbl>  
## 1  12.6
```

summarise() and group_by()

When paired with `group_by()`, it changes the unit of analysis from the whole data to specific subsets.

```
# Average departure delay by day
flights %>%
  group_by(year, month, day) %>%
  summarise(delay = mean(dep_delay, na.rm = TRUE))
```

```
## `summarise()` regrouping output by 'year', 'month' (override with `.groups` argument)
```

```
## # A tibble: 365 x 4
## # Groups:   year, month [12]
##   year month   day delay
##   <int> <int> <int> <dbl>
## 1  2013     1     1  11.5
## 2  2013     1     2  13.9
## 3  2013     1     3  11.0
## 4  2013     1     4   8.95
## 5  2013     1     5   5.73
## 6  2013     1     6   7.15
## 7  2013     1     7   5.42
## 8  2013     1     8   2.55
```

Useful summary functions

- Measures of location

- `mean()`, `median()`

- Measures of spread

- `sd()`, `IQR()`, `mad()`

- Measures of rank

- `min()`, `max()`, `quantile()`

- Measures of position

- `first()`, `nth()`, `last()`

- Counts

- `n()`, `sum(!is.na(x))`,
`n_distinct()`, `count()`

- Counts and proportions of logical values

- e.g. `sum(x > 10)`, `mean(y
= 0)`

Grouping by multiple variables

- Each summary peels off one level of the grouping.
- This works with sums and counts. **Be careful with rank-based statistics.**

```
# Daily flights
daily <- group_by(flights, year, month, day)
(per_day <- summarise(daily, flights = n()))

# Flights per month
(per_month <- summarise(per_day, flights = sum(flights)))

# Flights per year
(per_year <- summarise(per_month, flights = sum(flights)))
```

Ungrouping

- Removing grouping is easy, just use `ungroup()`.
- Useful when you need to do more data wrangling.

```
daily <- group_by(flights, year, month, day)

daily
```

```
## # A tibble: 336,776 x 19
## # Groups:   year, month, day [365]
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>         <int>
## 1  2013     1     1     517             515         2     830           819
## 2  2013     1     1     533             529         4     850           830
## 3  2013     1     1     542             540         2     923           850
## 4  2013     1     1     544             545        -1    1004          1022
## 5  2013     1     1     554             600        -6     812           837
## 6  2013     1     1     554             558        -4     740           728
## 7  2013     1     1     555             600        -5     913           854
## 8  2013     1     1     557             600        -3     709           723
## 9  2013     1     1     557             600        -3     838           846
## 10 2013     1     1     558             600        -2     753           745
## # ... with 336,766 more rows, and 11 more variables: arr_delay <dbl>,
```

Grouped mutates (and filters)

- The book suggests avoiding these, except in cases of quick manipulations.
- It's hard to check that the manipulation was done correctly.
- Book suggests checking out window functions `vignette("window-functions")`.

```
# Return only destinations with 365 flights
popular_dests <- flights %>%
  group_by(dest) %>%
  filter(n() > 365)

popular_dests
```


Questions/Discussion

- What examples/exercises did you find most useful from the reading?
- What examples/exercises gave you the most trouble?
- What is the most useful thing you took away from this chapter?